

Introduction to Regular Expressions

Cedric Arisdakessian

2022-09-02

A simple example

Your team compiled a metadata file for your project. One of the samples attribute is the location with the name of the island. However, each person has its own way to spell Hawaii:

- hawaii
- Hawaii
- Hawai'i
- Hawai'i

When you load your metadata into R, your island column has much more levels than it should have (25 instead of 5). It's because R doesn't know the different spellings correspond to the same island.

How would you fix this issue?

- 1 You don't have that many samples, you can just fix it manually.
- 2 You decide to hire an intern to fix it.
- 3 You look for all possible patterns in excel and replace them.
- 4 You have too many samples and you're not sure of all the spellings.
You decide to learn regular expressions.

What are regular expressions

It's a standardized language to define patterns of text, and it is commonly used to:

- Extract patterns
- Find/Replace patterns

Many tools use regex:

- Computer science languages (R, python, ... actually most of them)
- Linux command-line tools (grep, sed)
- Probably even excel in some way

By the end of this lecture, you should be able to decipher patterns like:

```
[A-Za-z0-9_\-]+\@[A-Za-z0-9_\-]+\.[a-z]{2,}$
```

Throughout the lecture, we are going to use R to see if a regex matches a text. We can do that with the `grep1(pattern, text)` function.

Before we start

- The regex language uses common characters to express patterns of text (ex: -, ?, .)
- What if you need to look for one of them in a text?
 - ⇒ By default, it will be assumed that you mean the regex interpretation.
 - ⇒ If you mean to match the actual symbol, you need to escape it (e.g. \. for a literal dot). In R, you will need 2 backslashes.

Dot symbol

The regex syntax is a set of symbols with a specific meaning:

- "." refers to any character.

```
grepl(".", "a")
```

```
## [1] TRUE
```

```
grepl("\\.", "a")
```

```
## [1] FALSE
```

```
grepl(".", "$")
```

```
## [1] TRUE
```

Matching specific location in the text

- "^" matches the beginning of the text.
- "\$" is similar to "^" but matches the end

```
# Look for letter "c" at the beginning of the text  
grepl("^c", c("cedric", "michaela"))
```

```
## [1] TRUE FALSE
```

```
# Look for letter "R" at the end of the text  
grepl(".R$", c("main.R", "Rstudio.app"))
```

```
## [1] TRUE FALSE
```

```
# bonus: what might be an issue with the previous example?
```

Combine regex

AND

If two regex follow each other (and are inside square brackets), the patterns combine:

```
grepl("a.a", "aba")
```

```
## [1] TRUE
```

```
grepl("a.a", "abba")
```

```
## [1] FALSE
```


OR

"x|y" will match text if it matches either x or y

```
grepl("a|b", "a")
```

```
## [1] TRUE
```

```
grepl("(a|b)c", "ac")
```

```
## [1] TRUE
```

Multiple choices and ranges

- The bracket "[xyz]" notation is similar and usually preferred for more than 2 choices
- The "[x-y]" notation refers to ranges (works for letters as well)

```
grepl("[ct]sv$", ".csv")
```

```
## [1] TRUE
```

```
grepl("[A-Z]", c("A", "a"))
```

```
## [1] TRUE FALSE
```

```
grepl("[A-Za-z0-9 ]", c("A", "a", "3", " "))
```

```
## [1] TRUE TRUE TRUE TRUE
```

Repetition of symbols

- "*" refers to 0 or more repetition of the previous pattern. Note that it's a bit different that its meaning as a glob pattern.
- "+" refers to 1 or more repetition of the previous pattern

```
grepl("a*", c("a", "", "aa"))
```

```
## [1] TRUE TRUE TRUE
```

```
grepl("a+", c("a", "", "aa"))
```

```
## [1] TRUE FALSE TRUE
```

```
grepl("aa*", c("a", "", "aa"))
```

```
## [1] TRUE FALSE TRUE
```

You can pick the number of repetitions with the "{n, m}" syntax

```
grepl("ba{2,3}b", c("baab", "baaaab"))
```

```
## [1] TRUE FALSE
```

```
grepl("a{2,}", c("a", "aa", "aaaa"))
```

```
## [1] FALSE TRUE TRUE
```

Warning

```
grepl("a{,2}", "aaaaa")
```

```
## [1] TRUE
```

```
grepl("^a{,2}$", "aaaaa")
```

```
## [1] FALSE
```

Negation

When at the beginning of square brackets, the "^" symbol negates a set of characters:

```
grepl("b[^bB]b", c("bbb", "bBb", "bab"))
```

```
## [1] FALSE FALSE  TRUE
```

Summary

- `"."`: any character.
- `x|y`: x OR y
- `[xyz]`: x OR y OR z
- `[x-y]`: any character in range (numbers or letters)
- `[^xyz]`: anything but x, y or z
- `"^"`: matches the beginning of the text.
- `"$"`: matches the end of the text
- `"*"`: 0+ repetitions
- `"+"`: 1+ repetitions
- `{n,m}`: between n and m repetitions

More regex

There's actually more to the regex syntax that is presented here. Once you're familiar with the ones presented here, it might be interesting looking into:

- capture groups (select one or multiple groups and reorganize them)
- positive/negative lookahead (need to understand capture groups first)

Using regex in R

- `grep`, `grepl` : Find in string
- `gsub` : string substitution

```
# gsub(pattern, replacement, text)  
gsub("a", "b", "aaa")
```

```
## [1] "bbb"
```

```
gsub("a", "b", c("aaa", "cac"))
```

```
## [1] "bbb" "cbc"
```


Using regex in linux: grep

grep [PATTERN] [FILE] find the lines matching a given pattern in a file. (many additional options)

```
# Count the number of sequences in a fasta file
$ grep -c "^>" sequences.fasta
```

```
# Count the number of sequences in a fastq file
$ grep -c "^@" reads.fastq # roughly
$ grep -c "^@M01" reads.fastq # a bit better
```

```
# Retrieve the entries labelled as S.aureus in the header of a f
# (needs to be 2-line formatted)
$ grep -A1 ">.*S\.aureus.*"
# or
$ grep -A1 "S\.aureus"
```

Using regex in linux: sed

Very complex command. But basic functionalities are more manageable:

- `sed 's/[pattern]/[replacement]/'` (replace the first occurrence of [pattern] with [replacement] in each line)
- `sed 's/[pattern]/[replacement]/g'` (replace any occurrence of [pattern] with [replacement] in each line)
- `sed '[line_pattern]/s/[pattern]/[replacement]/g'` (match only lines with pattern)
- `sed '[line_pattern]/!s/[pattern]/[replacement]/g'` (negate line pattern)

And much more.

```
# Replace excel "#VALUE" in csv file (inplace)  
$ sed -i 's/#VALUE/NaN/g' metadata.csv  
# More complex: Remove all N nucleotides in a fasta file  
$ sed '^>/!s/N//g' sequences.fasta > sequences_N-removed.fasta
```